

Tutorial: Identifying Software Triviality via Fine-grained and Dataflow-based Value Profiling



Zhibo Xuan
Beihang University
Hands-on Tutorial @ CLUSTER24



北京航空航天大學
BEIHANG UNIVERSITY

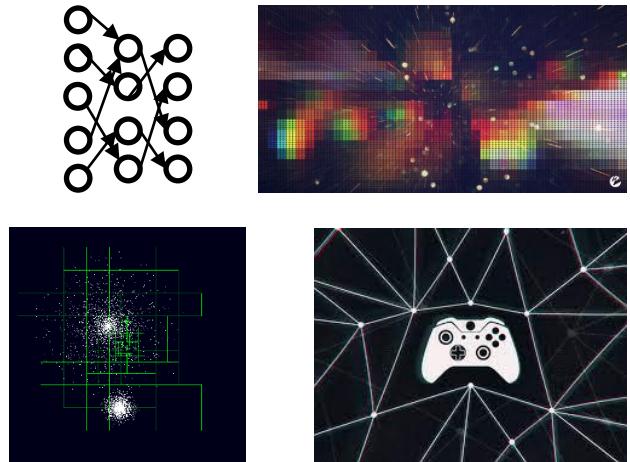
Outline

- Introduction & Background
- Understanding Software Triviality
- Dataflow-based Triviality Detection
- Evaluation
- Hands-on Tutorial
 - Installation
 - Case Study – Backprop
 - Case Study – IS Benchmark
 - Case Study – fotonik3d

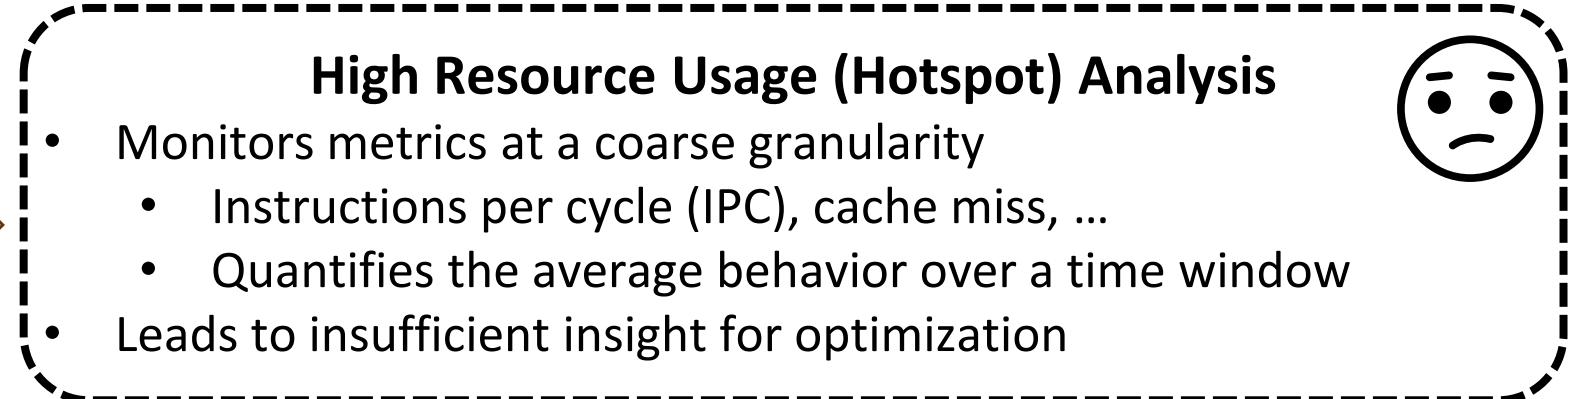
Outline

- Introduction & Background
- Understanding Software Triviality
- Dataflow-based Triviality Detection
- Evaluation
- Hands-on Tutorial
 - Installation
 - Case Study – Backprop
 - Case Study – IS Benchmark
 - Case Study – fotonik3d

From Resource Usage to Resource Wastage



Profiling

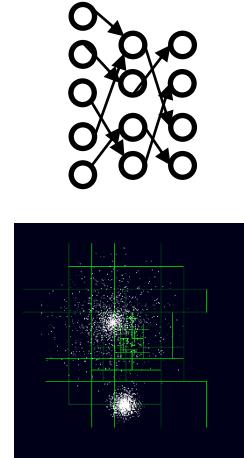


Code regions with high resource usage are likely to reveal optimization potentials, it still cannot reflect whether the resources utilized effectively.

- *log(1) in a loop is wasteful use of FPU, although it has high resource usage as well as high IPC.*

**Performance analysis needs to focus on resource wastage
in addition to resource usage**

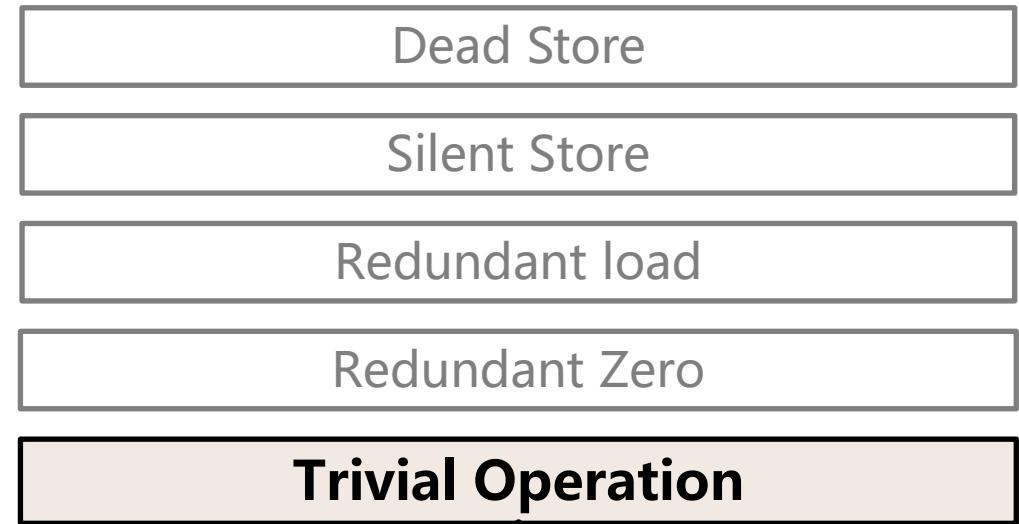
Resource Wastage - Trivial Operation



*sophisticated control flows and
deep layers of abstractions*

Unexpected Software
Inefficiencies

*e.g., 0^*x , $\log(0)$, ... in loops*



**Pervasively exist in real-world applications with
notable negative performance impacts**

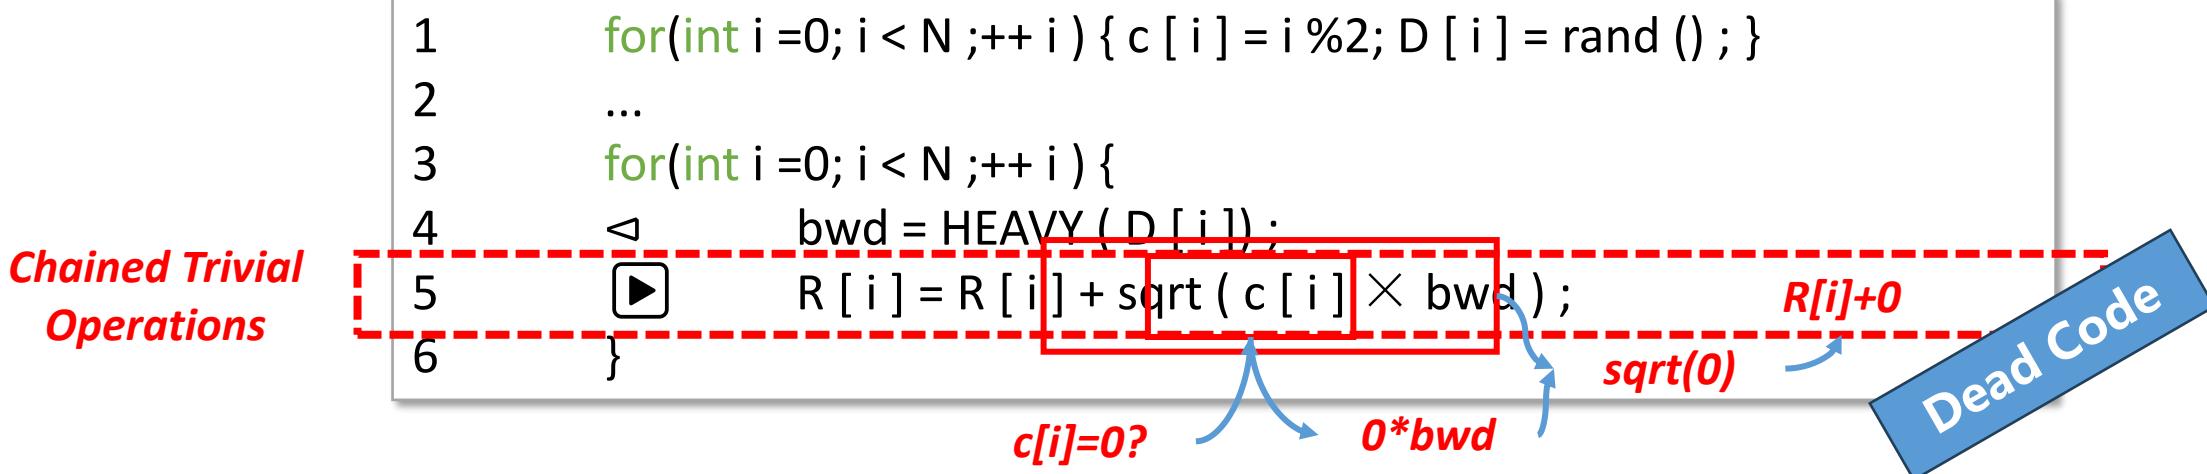
- A trivial operation will always result in **the same value when specific conditions are satisfied**.
- Executing trivial operations leads to a **waste of functional units and memory bandwidth**, revealing new performance optimization opportunities.

Trivial Operation – An example

```
1   for(int i =0; i < N ;++ i ) { c [ i ] = i %2; D [ i ] = rand () ; }
2   ...
3   for(int i =0; i < N ;++ i ) {
4       ◀      bwd = HEAVY ( D [ i ] );
5       ▶      R [ i ] = R [ i ] + sqrt ( c [ i ] × bwd ) ; Trivial Operations
6   }
```

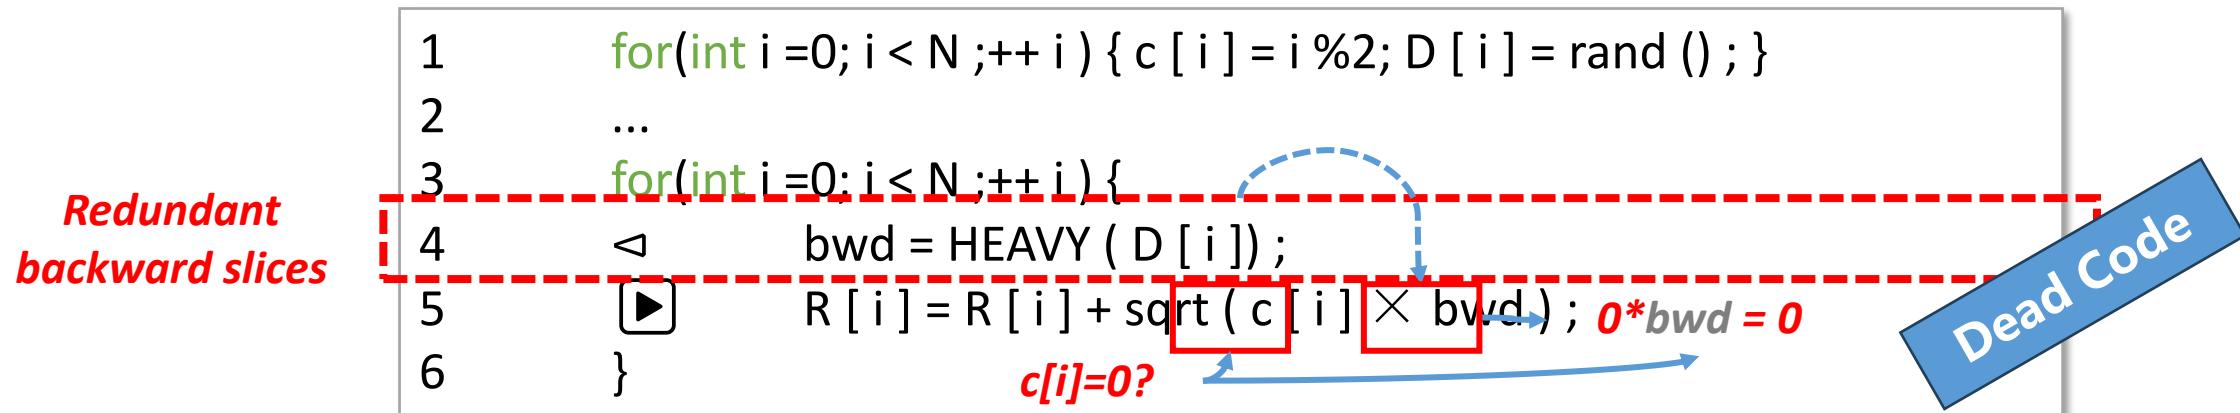
- ***Chained trivial operations*** are trivial operations that can be triggered in sequence with the same conditions
- ***Redundant backward slices*** are the dead codes when the trivial operations are eliminated with the specific conditions

Trivial Operation – An example



- ***Chained trivial operations*** are trivial operations that can be triggered in sequence with the same conditions
- ***Redundant backward slices*** are the dead codes when the trivial operations are eliminated with the specific conditions

Trivial Operation – An example



- ***Chained trivial operations*** are trivial operations that can be triggered in sequence with the same conditions
- ***Redundant backward slices*** are the dead codes when the trivial operations are eliminated with the specific conditions

Trivial Operation – An example

```
1   for(int i =0; i < N ;++ i ) { c [ i ] = i %2; D [ i ] = rand () ; }
2   ...
3   for(int i =0; i < N ;++ i ) {
4     ▲ bwd = HEAVY ( D [ i ] );
5     ▶ R [ i ] = R [ i ] + sqrt ( c [ i ] × bwd ) ;
6   }
```

Optimized

```
3   for(int i =0; i < N ;++ i ) {
4     if (c[i]!=0) {
5       ▲ bwd = HEAVY ( D [ i ] );
6       ▶ R [ i ] = R [ i ] + sqrt ( c [ i ] × bwd ) ;
7     }
8   }
```

Optimizes performance by avoiding redundant computations and memory accesses.

- ***Chained trivial operations*** are trivial operations that can be triggered in sequence with the same conditions
- ***Redundant backward slices*** are the dead codes when the trivial operations are eliminated with the specific conditions

Outline

- Introduction & Background
- Understanding Software Triviality
- Dataflow-based Triviality Detection
- Evaluation
- Hands-on Tutorial
 - Installation
 - Case Study – Backprop
 - Case Study – IS Benchmark
 - Case Study – fotonik3d

Software Trivialities

■ **Absorbing Triviality**

- The triviality is absorbing if the trivial condition makes other operand irrelevant to the result when satisfied.

■ **Identical Triviality**

- The triviality is identical if the operation result equals to other operand when the trivial condition is satisfied.

■ **Functional Triviality**

- The triviality is functional if the trivial condition $x \equiv c$ leads to a constant operation result where $c \in \{0, 1, F\}$.

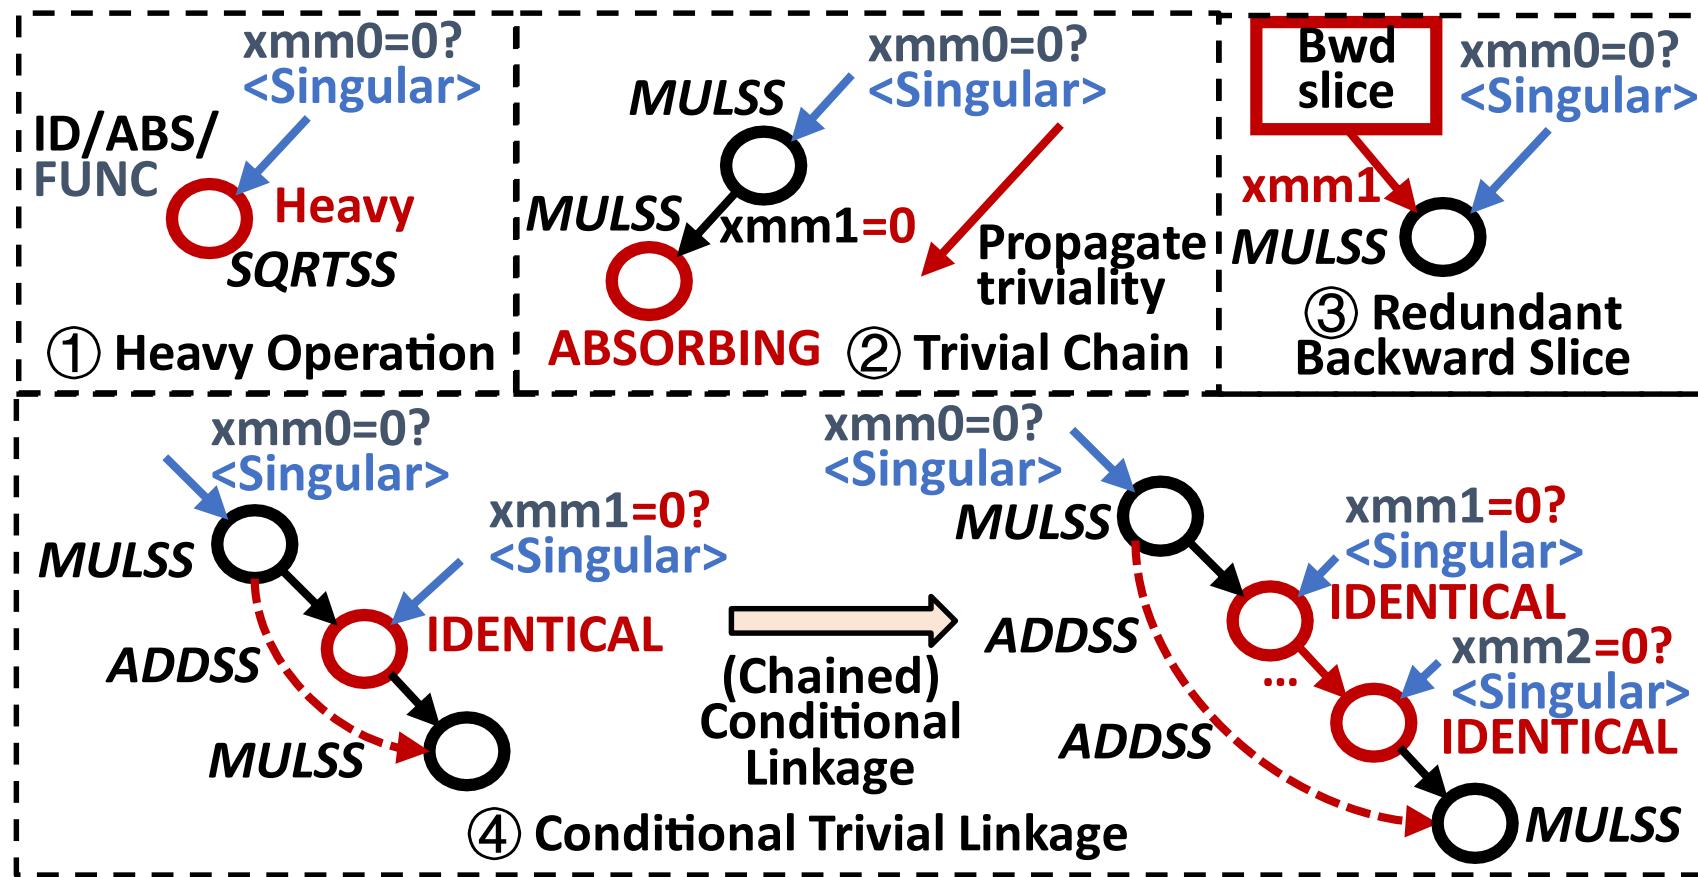
An operation is an instruction (e.g., MULSS, SQRTSS) or a *wrapped math function* (e.g., *exp*, *log*).

Operation	Trivial Condition	Type	Results
A & B	A=F/B=F	Identical Absorbing	B/A 0
	A=0/B=0		
A B	A=F/B=F	Absorbing Identical	F B/A
	A=0/B=0		
A + B	A=0/B=0	Identical	B/A
	B=0		
A - B	A=0/B=0	Identical	A
	A=1/B=1		
A * B	A=0/B=0	Absorbing Identical	0 B/A
	A=1/B=1		
A / B	A=0	Absorbing Identical	0 A
	B=1		
sqrt(A)	A=0	Functional	0 1
	A=1		

* *F* indicates a full trivial condition, in which all bits are set to 1 with the specified data length (e.g., *0xff* for *int8*)

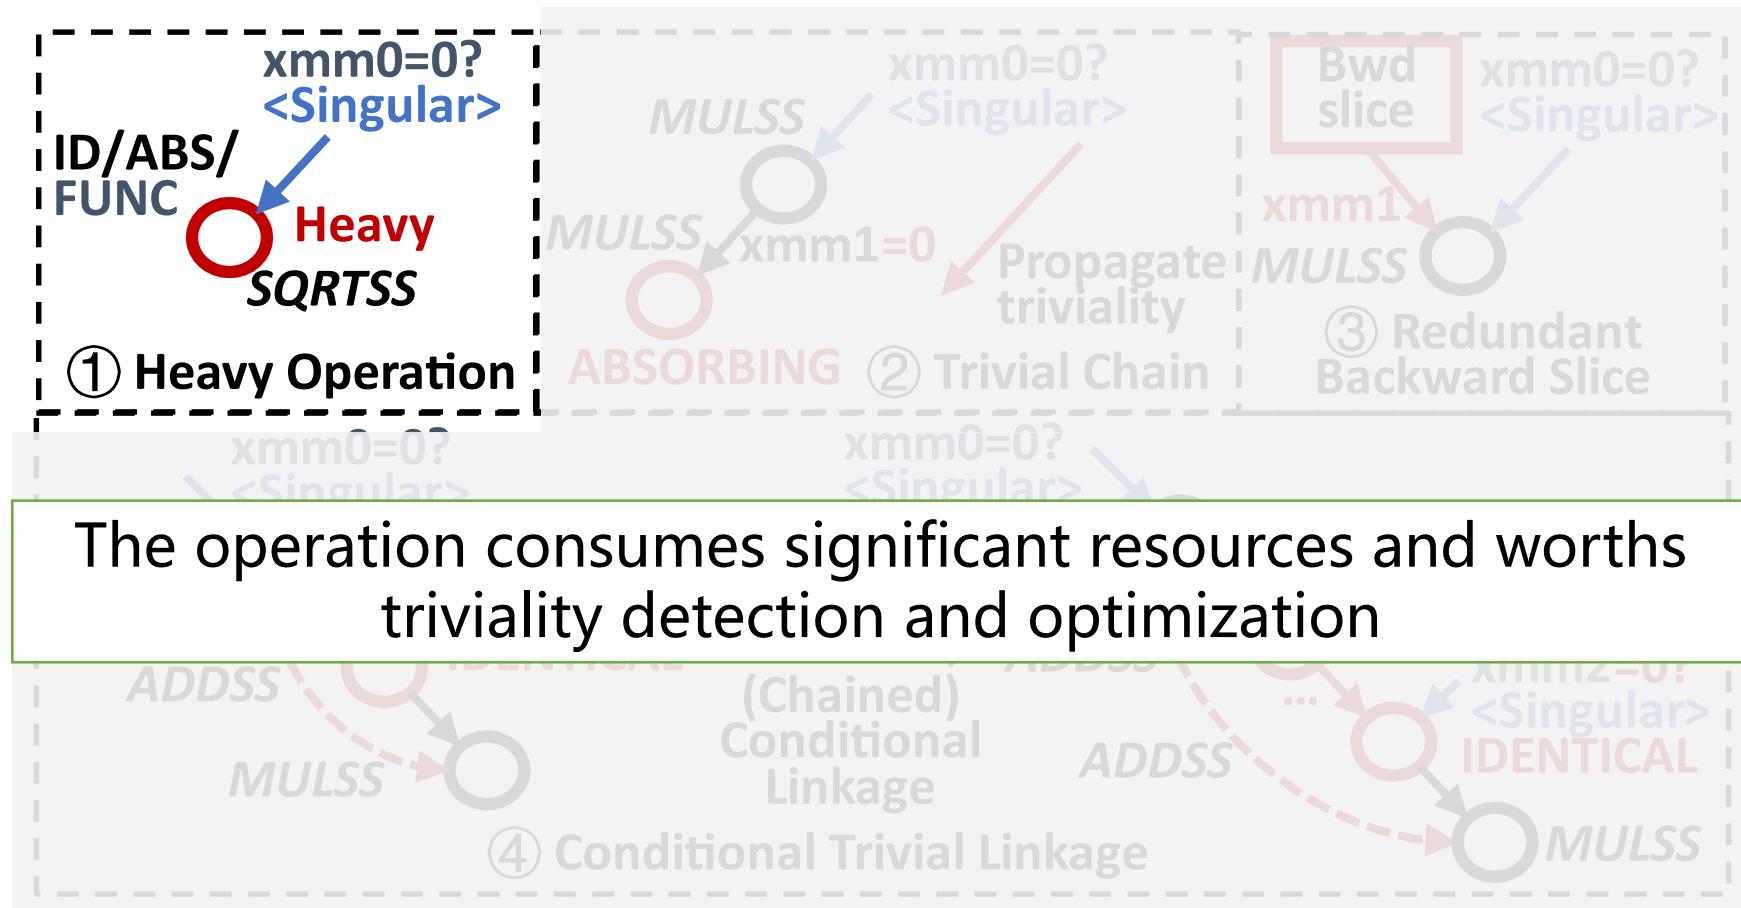
Singular Trivial Condition

- A ***Singular Trivial Condition (STC)*** is a precondition on a specific operand that triggers software trivialities with a significant performance impact.



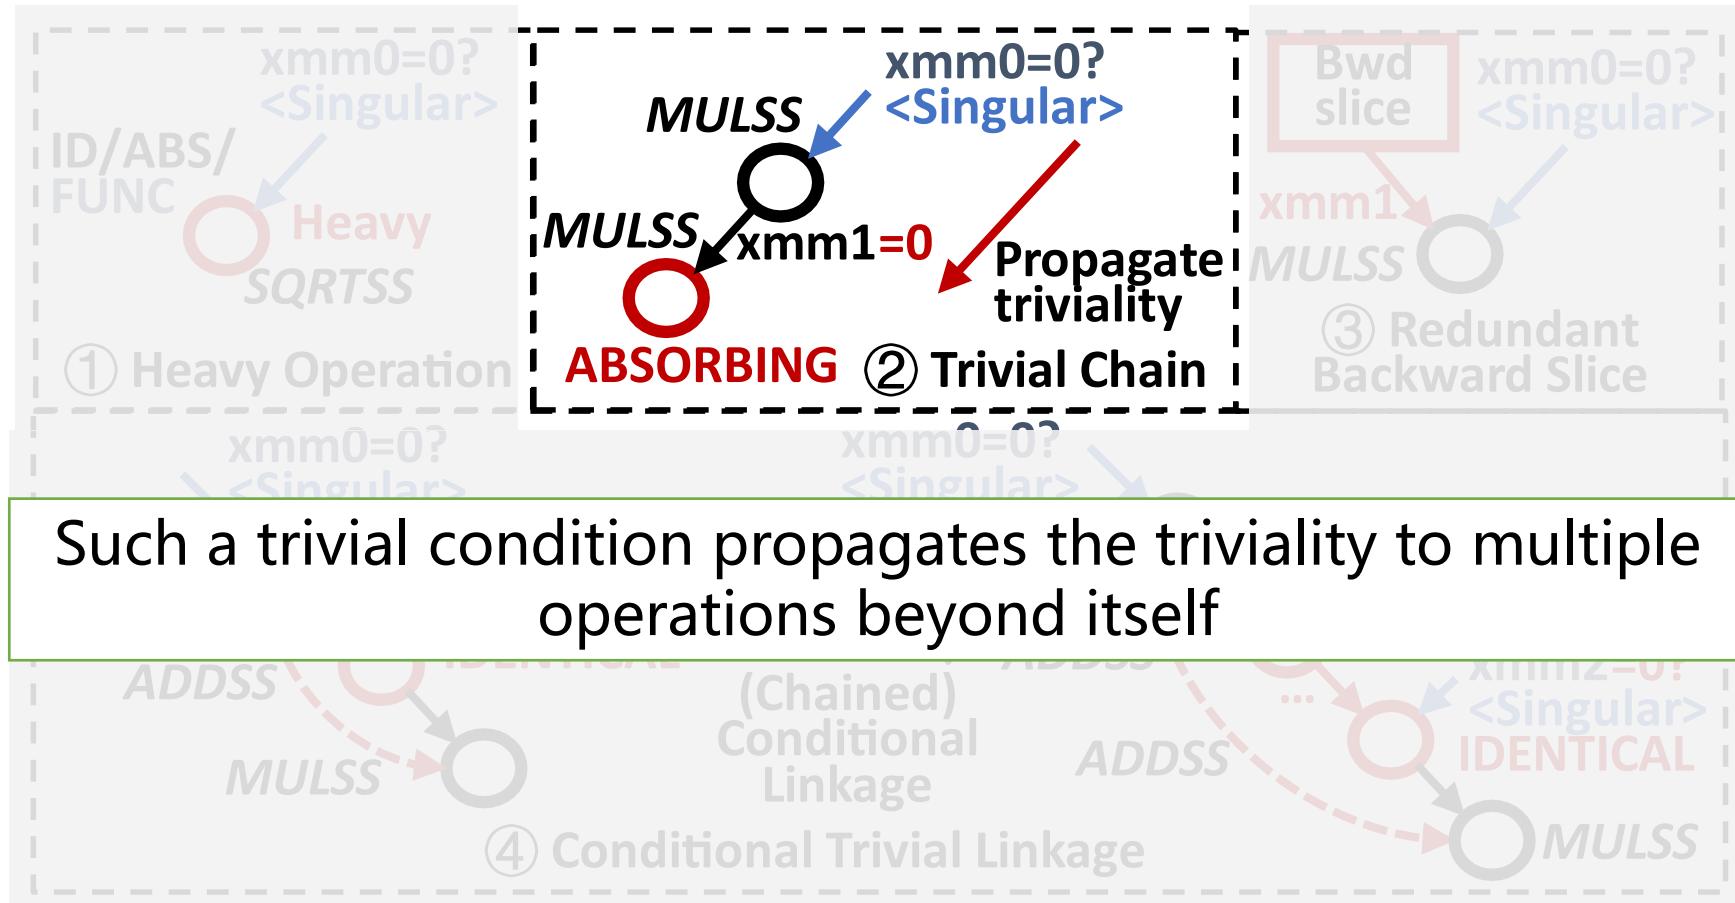
Singular Trivial Condition

- A ***Singular Trivial Condition (STC)*** is a precondition on a specific operand that triggers software trivialities with a significant performance impact.



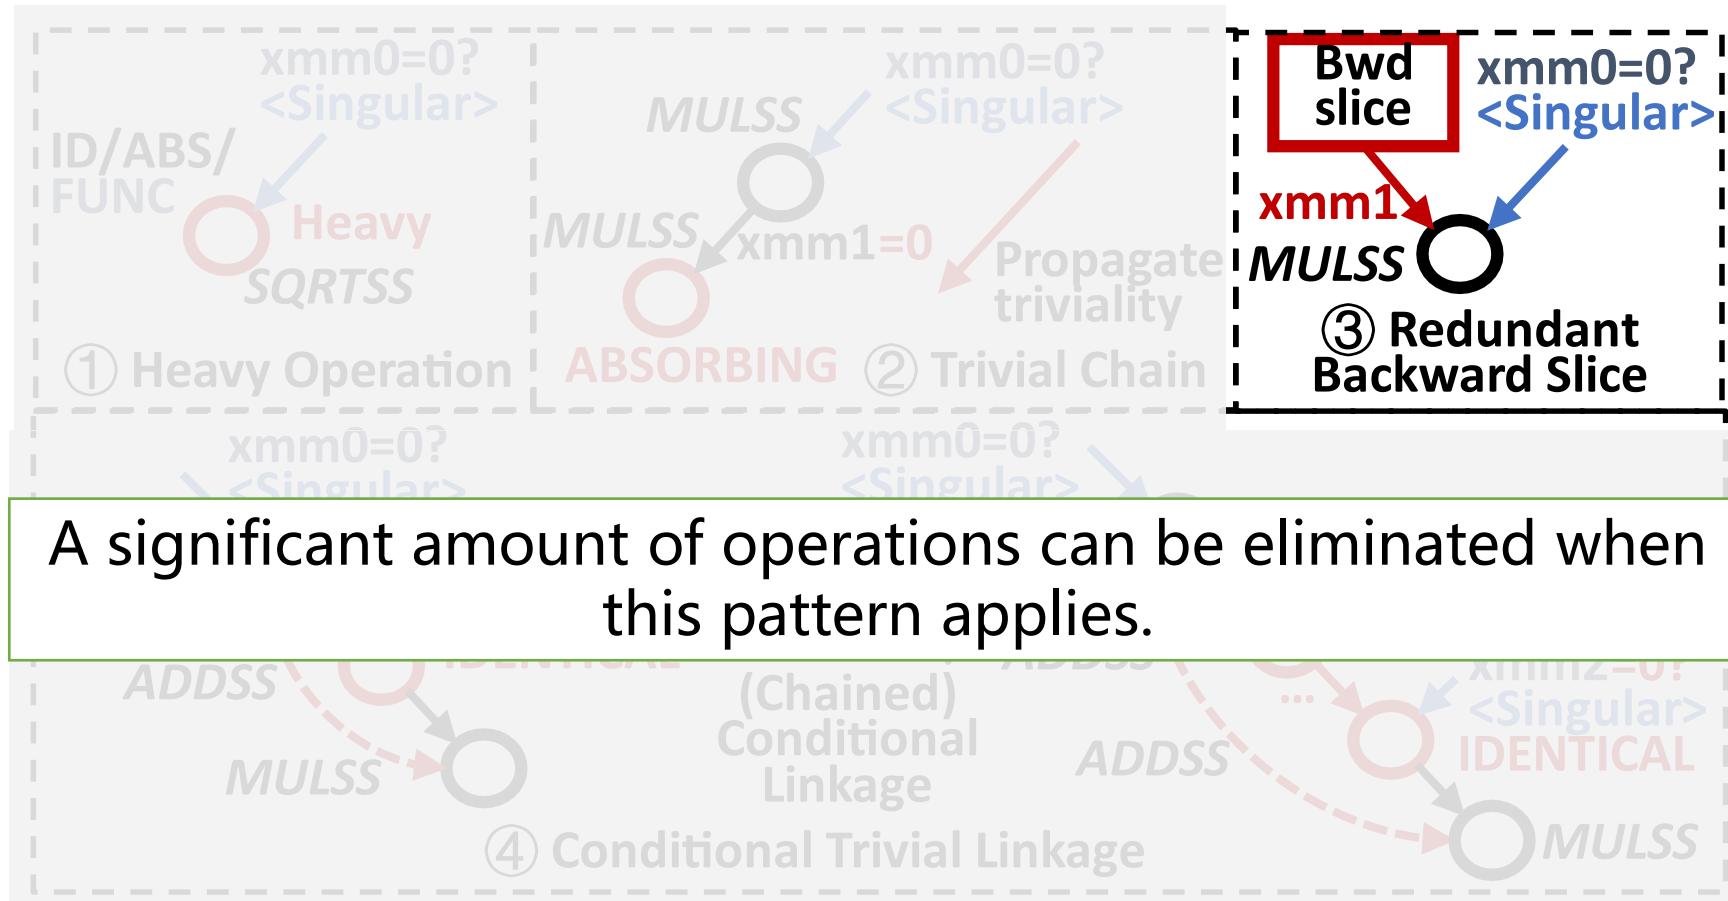
Singular Trivial Condition

- A ***Singular Trivial Condition (STC)*** is a precondition on a specific operand that triggers software trivialities with a significant performance impact.



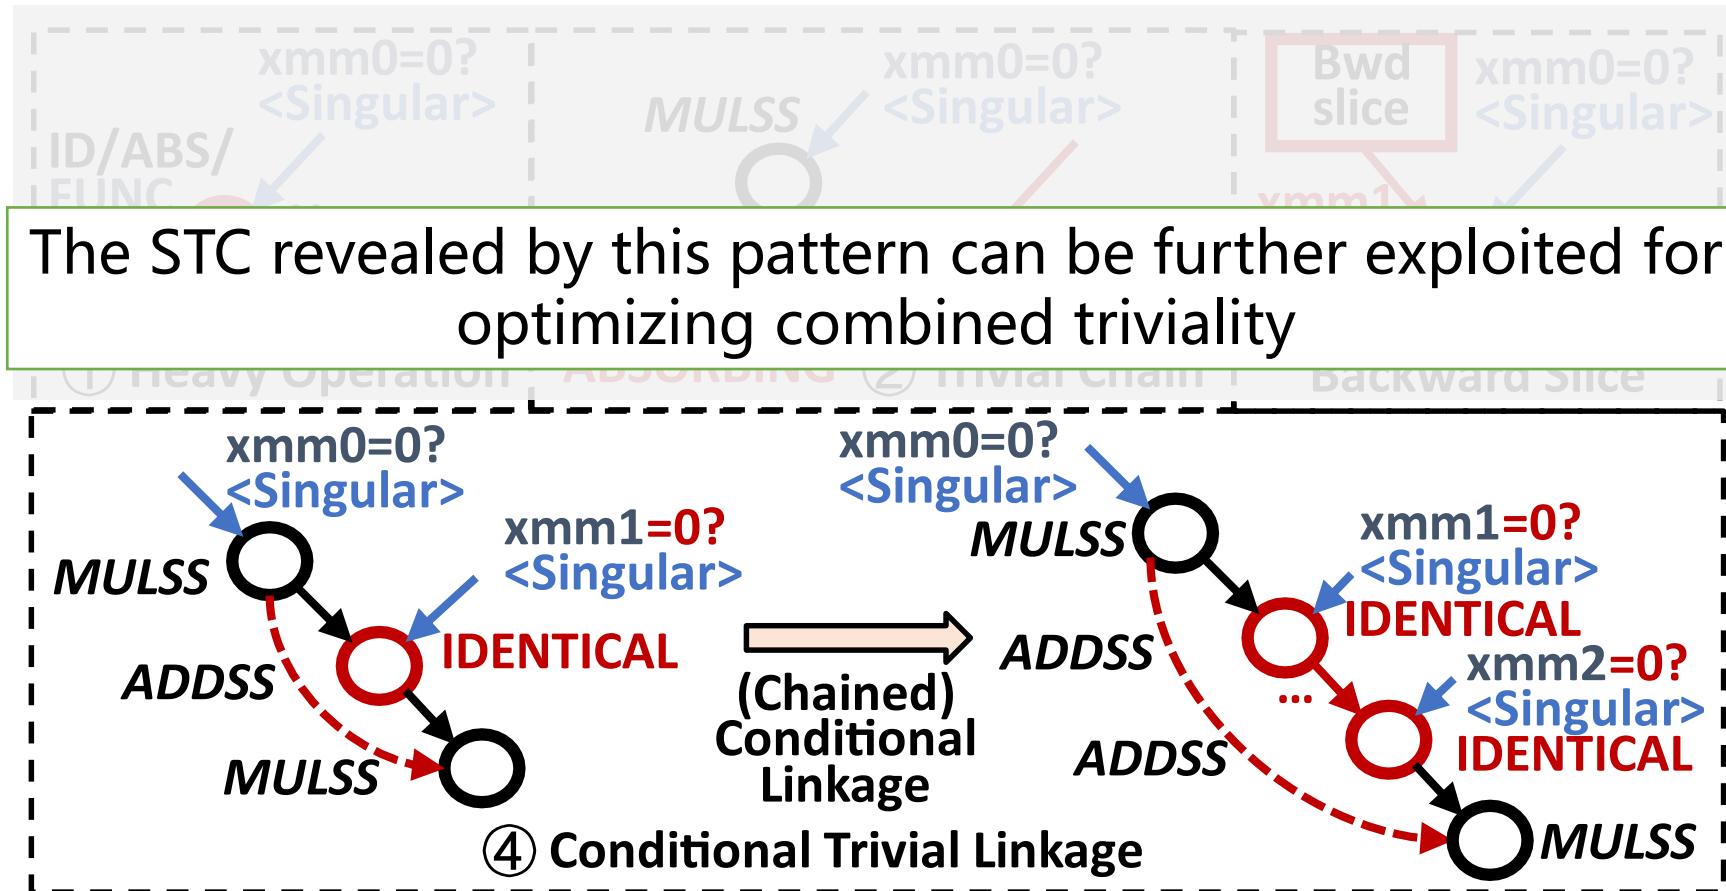
Singular Trivial Condition

- A ***Singular Trivial Condition (STC)*** is a precondition on a specific operand that triggers software trivialities with a significant performance impact.



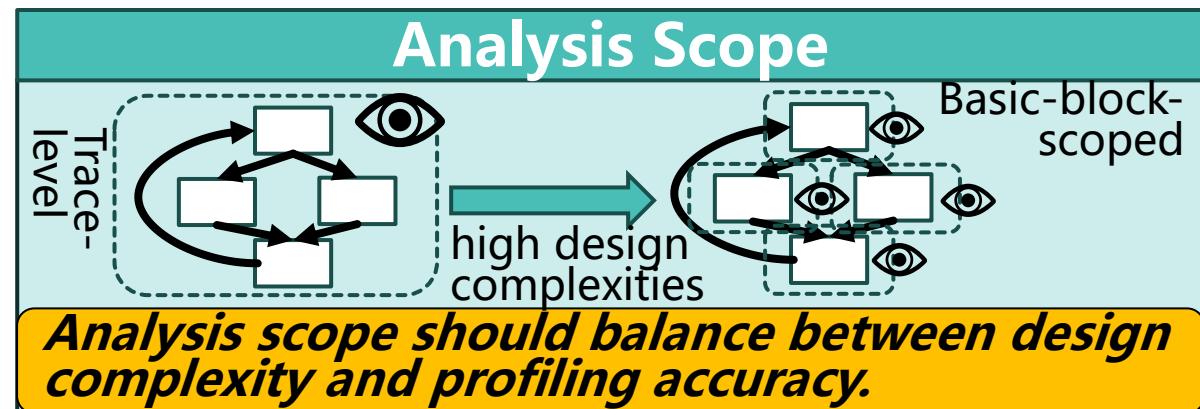
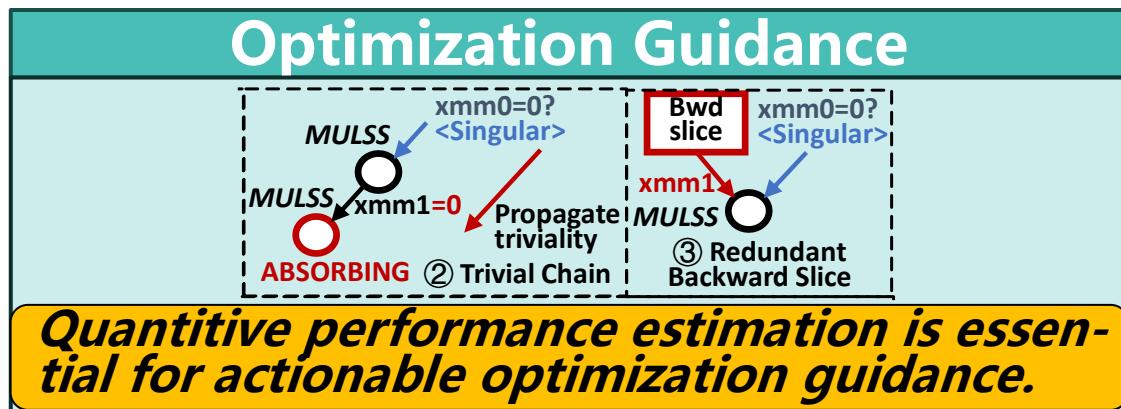
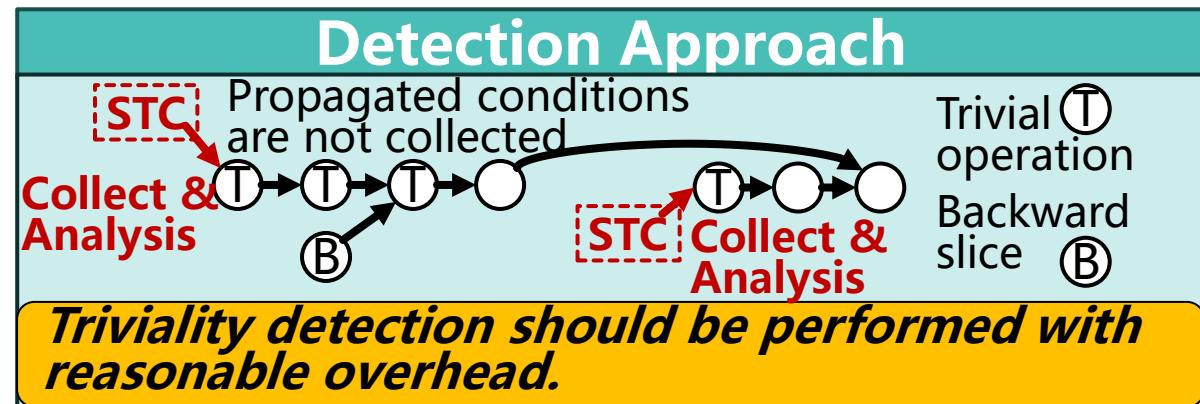
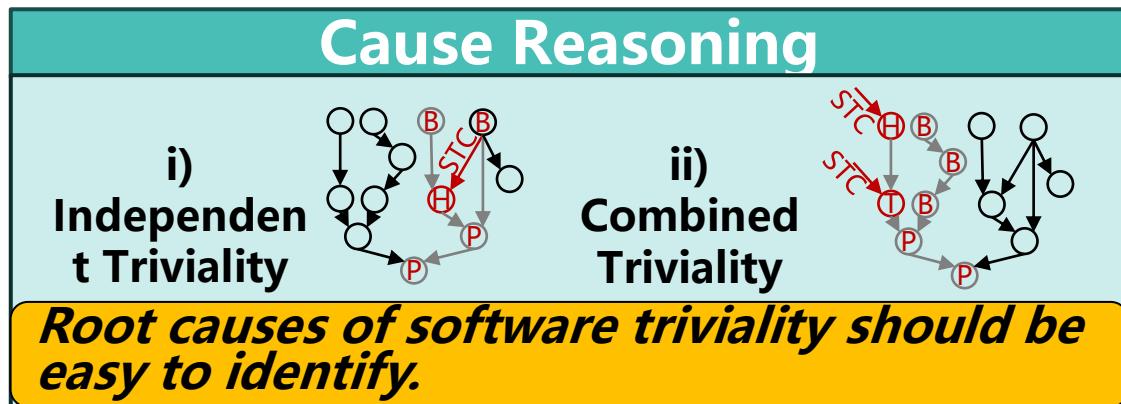
Singular Trivial Condition

- A ***Singular Trivial Condition (STC)*** is a precondition on a specific operand that triggers software trivialities with a significant performance impact.



Identifying the Causes

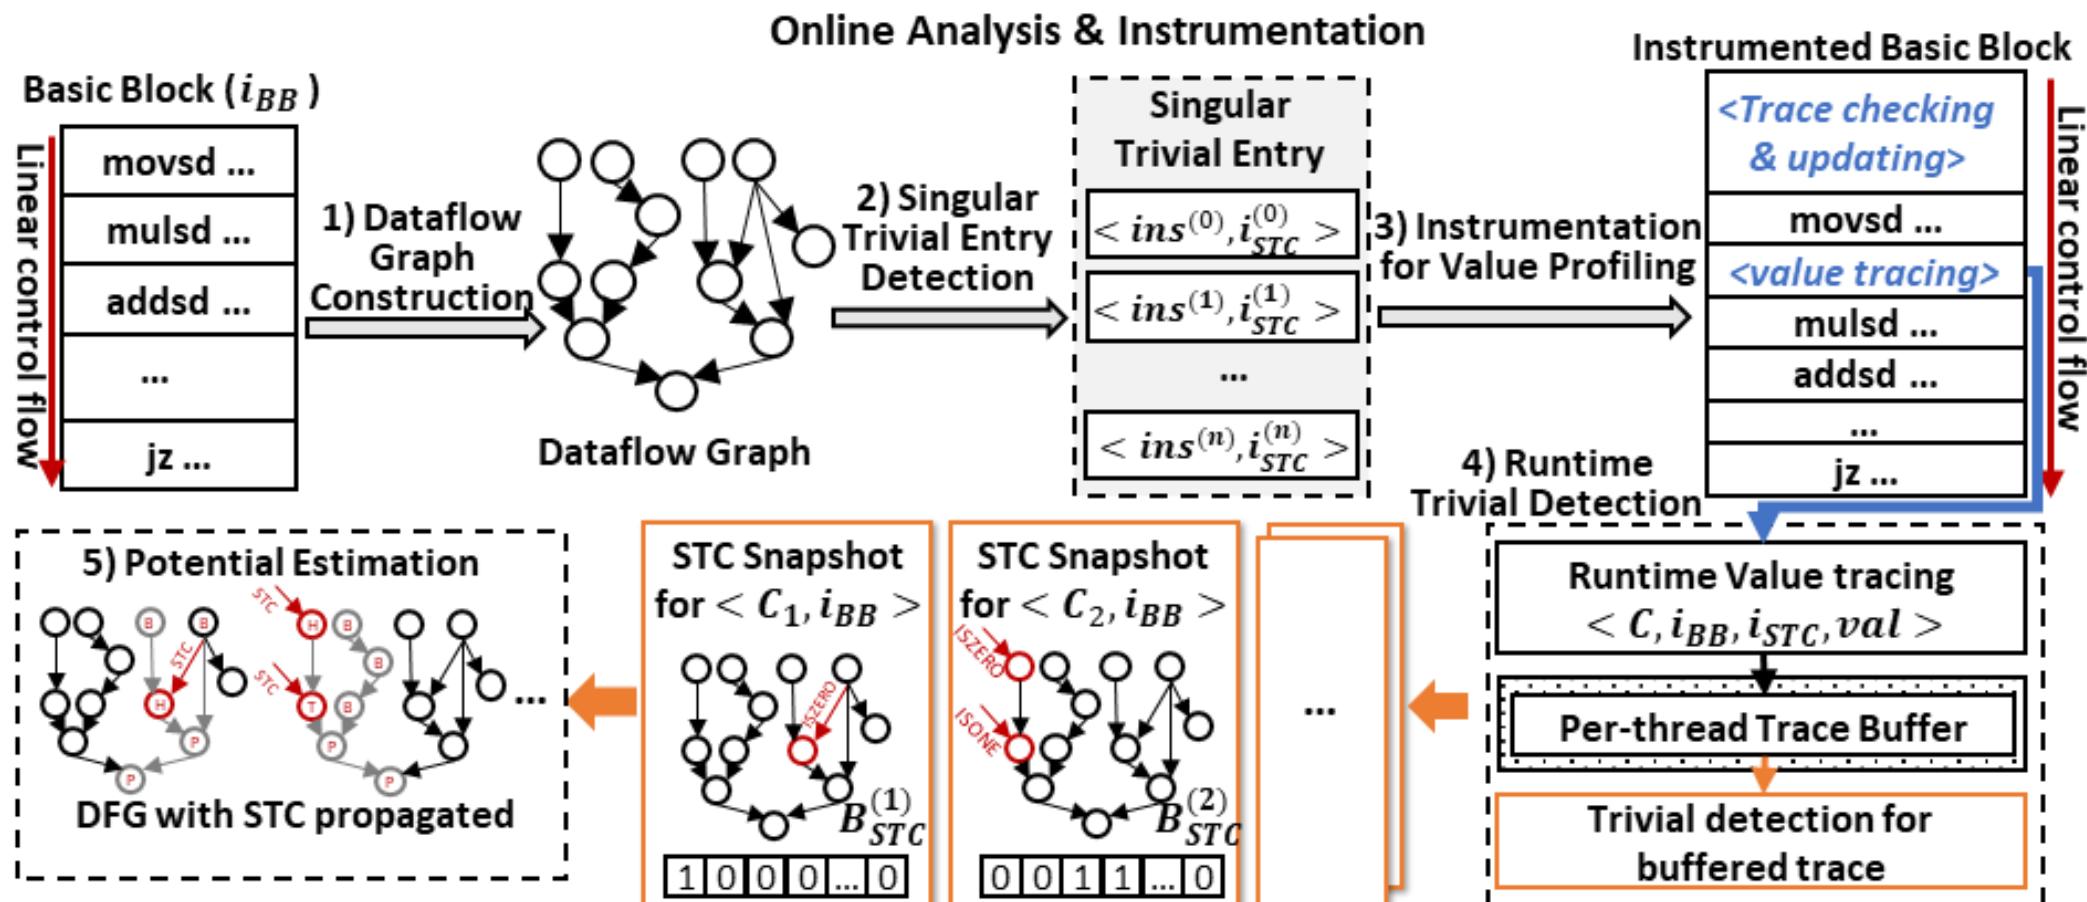
- We can derive four principles based on the above observed patterns to identify the root causes of software trivialities within acceptable overhead.



Outline

- Introduction & Background
- Understanding Software Triviality
- Dataflow-based Triviality Detection
- Evaluation
- Hands-on Tutorial
 - Case Study – Backprop
 - Case Study – IS Benchmark
 - Case Study – fotonik3d

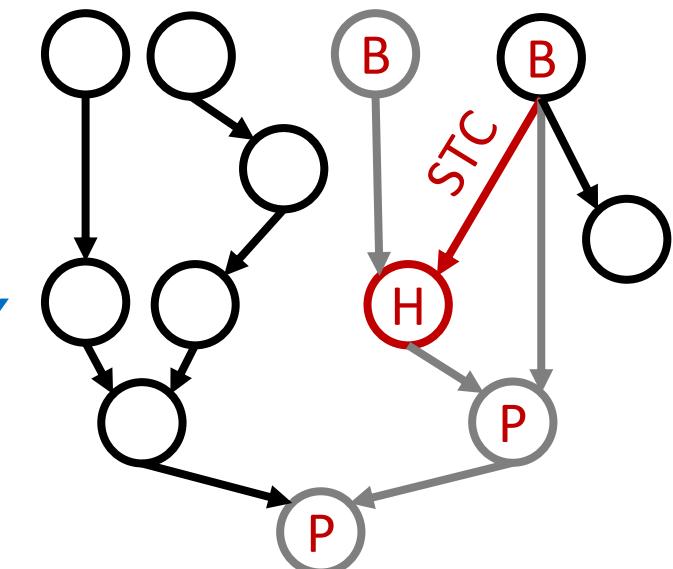
Dataflow-based Triviality Detection – Overview



- We develop a fine-grained dataflow-based value profiler *TrivialSpy* to detect and estimate the optimization potential of software triviality.

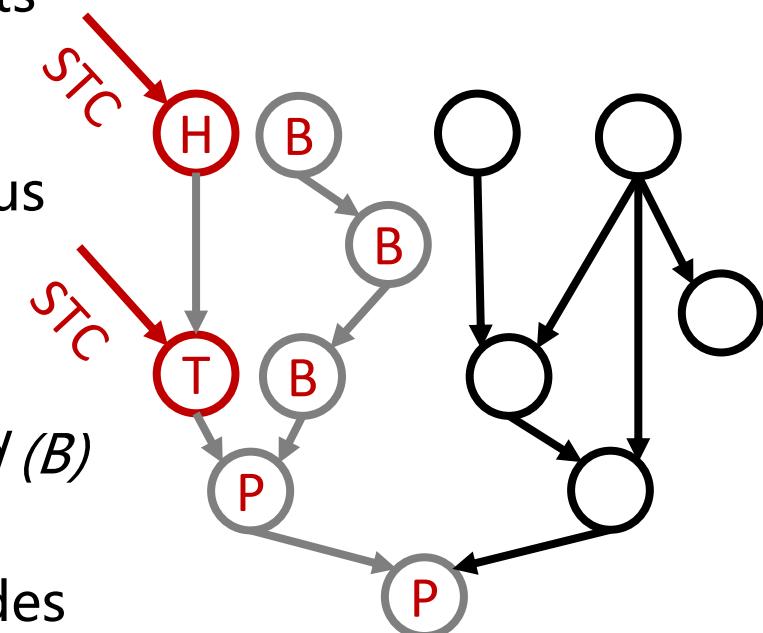
Singular Trivial Entry Detection – Independent Triviality

- **Trivial condition selection**
 - Select first *non-propagated* trivial condition in BFS
- **Pattern detection**
 - Applying **Pattern 1~3**
 - *heavy operation, trivial chain, redundant backward slice*
 - Mark the nodes as *heavy (H)*, *propagated (P)*, and *backward (B)*
 - **detect and mark all root causes of specific software triviality**
- **Dead code detection**
 - all *heavy* and *propagated* nodes will be marked as dead codes
 - *backward* is considered as dead codes *iff. all its children are detected as dead codes*

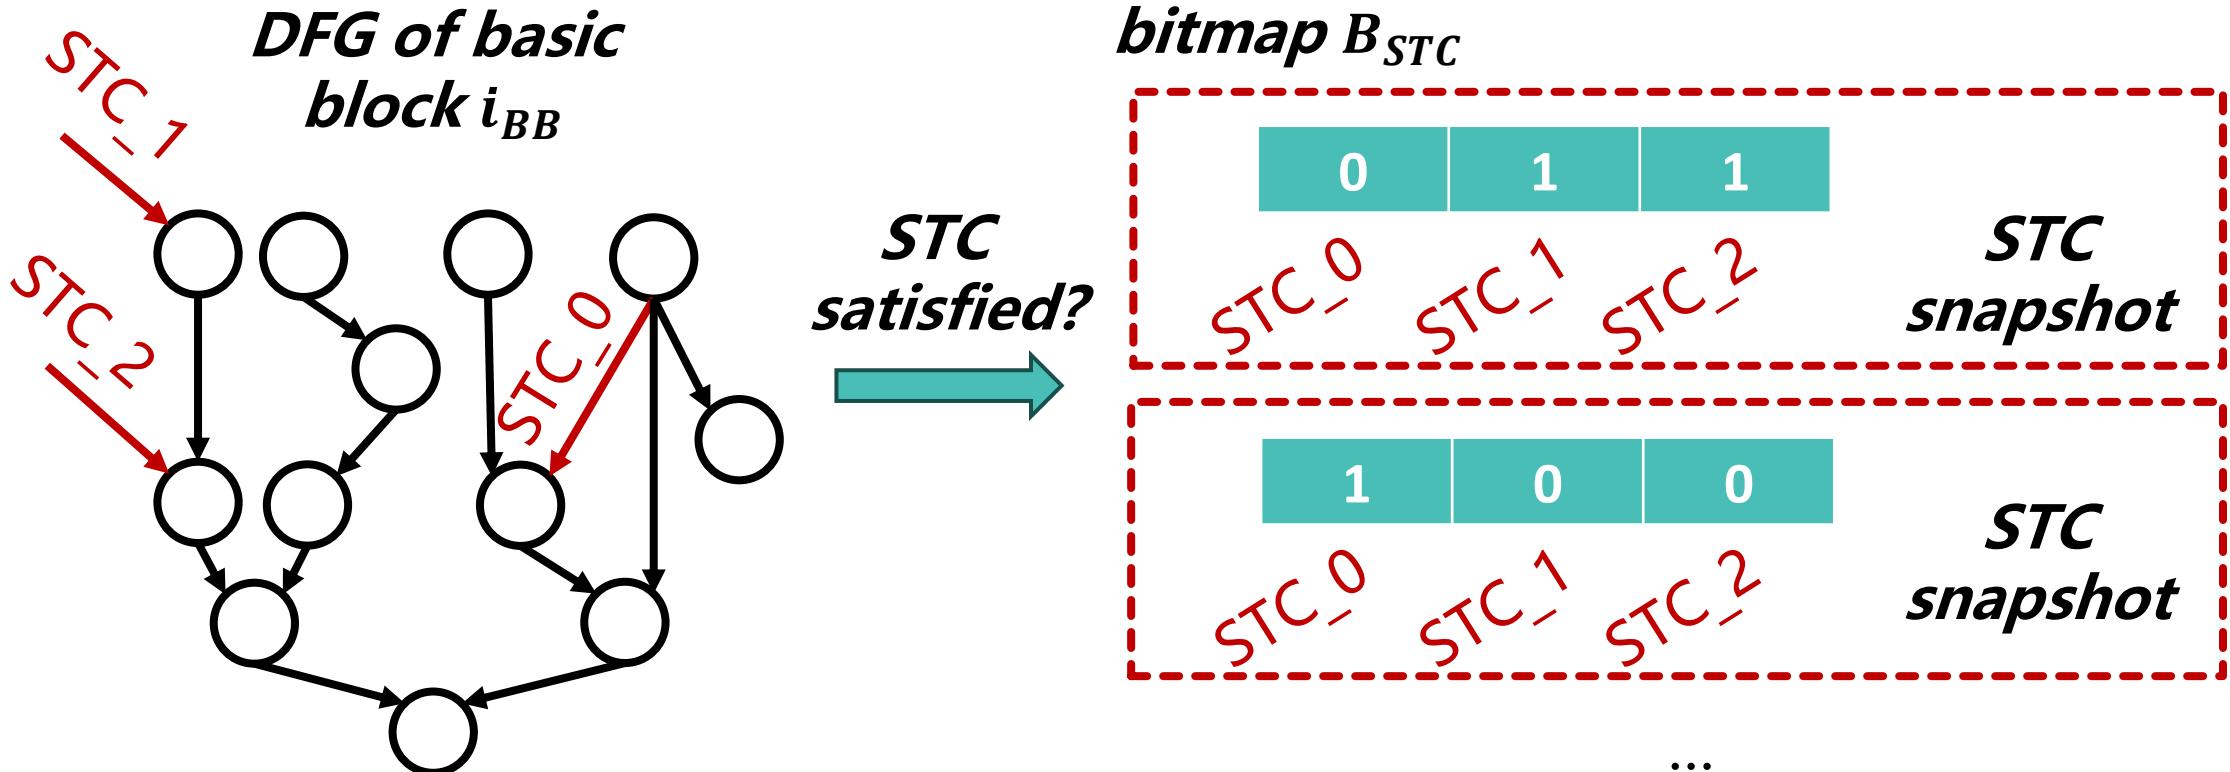


Singular Trivial Entry Detection – Combined Triviality

- **Conditional linkage discovery**
 - Scans the nodes in the DFG with *identical triviality*
 - Each has a trivial parent generating results to make one of its children trivial
 - **Backward reasoning**
 - Analyze if the trivial conditions are propagated from previous trivial operations in the DFG
 - **Pattern detection**
 - Applying **Pattern 1~3**
 - Mark the nodes as *heavy (H)*, *propagated (P)*, and *backward (B)*
 - **Dead code detection**
 - all *heavy* and *propagated* nodes will be marked as dead codes
 - *backward* is considered as dead codes iff. *all its children are detected as dead codes*



Runtime Trivial Detection

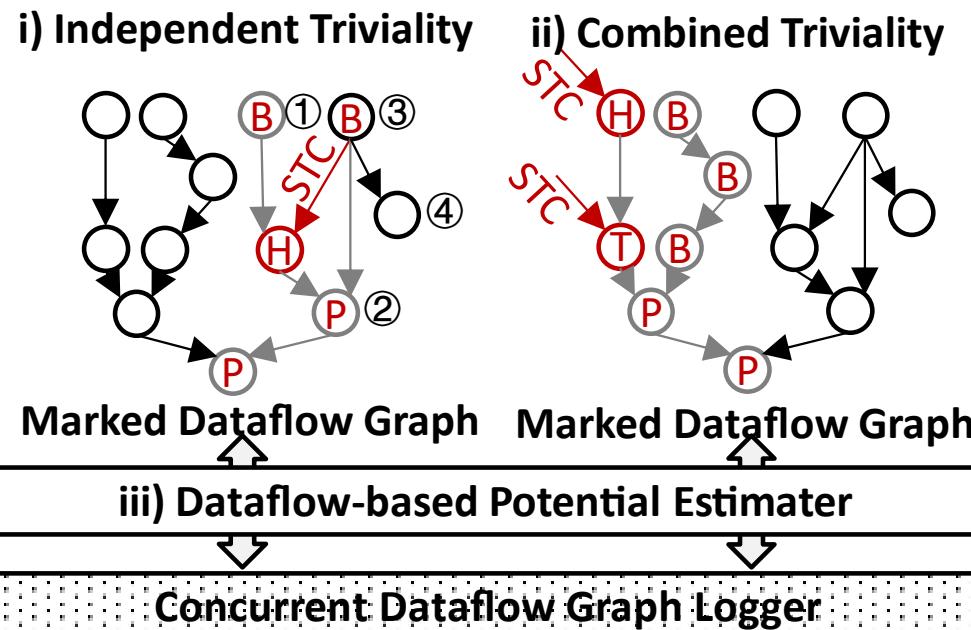


Accumulates the STC snapshots for the total number of execution with the parent' s calling context of the basic block

Potential Estimation

- **Expected Benefit (EB) & Branching Benefit (BB)**

- We mark the corresponding DFG by constantly propagating with the detected STC recorded in the snapshot (similar to *singular trivial entry detection*)
- The performance potential of a trivial condition is estimated as the accumulated cost of operations marked as dead by the previous propagating phase



$$Cost(V) = Latency(V) + Latency_{mem} \times N_{mem}$$

$$Cost(G) = \sum_{V \in G} Cost(V)$$

$$EB = N_{trivial} \times (Cost(G) - Cost(G_{fast}))$$

$$BB = EB - N \times Cost(C)$$

BB indicates the expected performance improvement of branch optimization

Reporting software trivialities

Branching Benefit: $BR_{local} (BB_{local}/BB_{total})$
Importance: $BI_{local} (BB_{local}/Cost_{total})$

1

Performance potential

Expected Benefit: $ER_{local} (EB_{local}/EB_{total})$
Importance: $EI_{local} (EB_{local}/Cost_{total})$

2

^^ Trivial Ratio: $R_{local} (N_{trivial}/N_{total})$ ^^

3

+++ DFG Caller CCT Info +++
<instruction>@<func>[<file>:<line>]
...

4

Calling context

===== DFGLog from Thread *tid* =====
exe count: *N*
... (profiled results of $Cost(G)$, EB , HIR , TCR , $RBSR$, CTR)

5

Detailed metric

+++++ Singular Trivial Condition(s):
<i> <src, dst>: <opnd>, <STC val>, <isSingular>
...

6

STC

==> detailed node info:
[i] <instruction>@<func>[<file>:<line>] <[B][P][H]>
...

7

==> detailed edge info:
<i> <src, dst>: <opnd>, <propagated value>, <isSingular>
...

DFG with source lines

Outline

- Introduction & Background
- Understanding Software Triviality
- Dataflow-based Triviality Detection
- Evaluation
- Hands-on Tutorial
 - Installation
 - Case Study – Backprop
 - Case Study – IS Benchmark
 - Case Study – fotonik3d

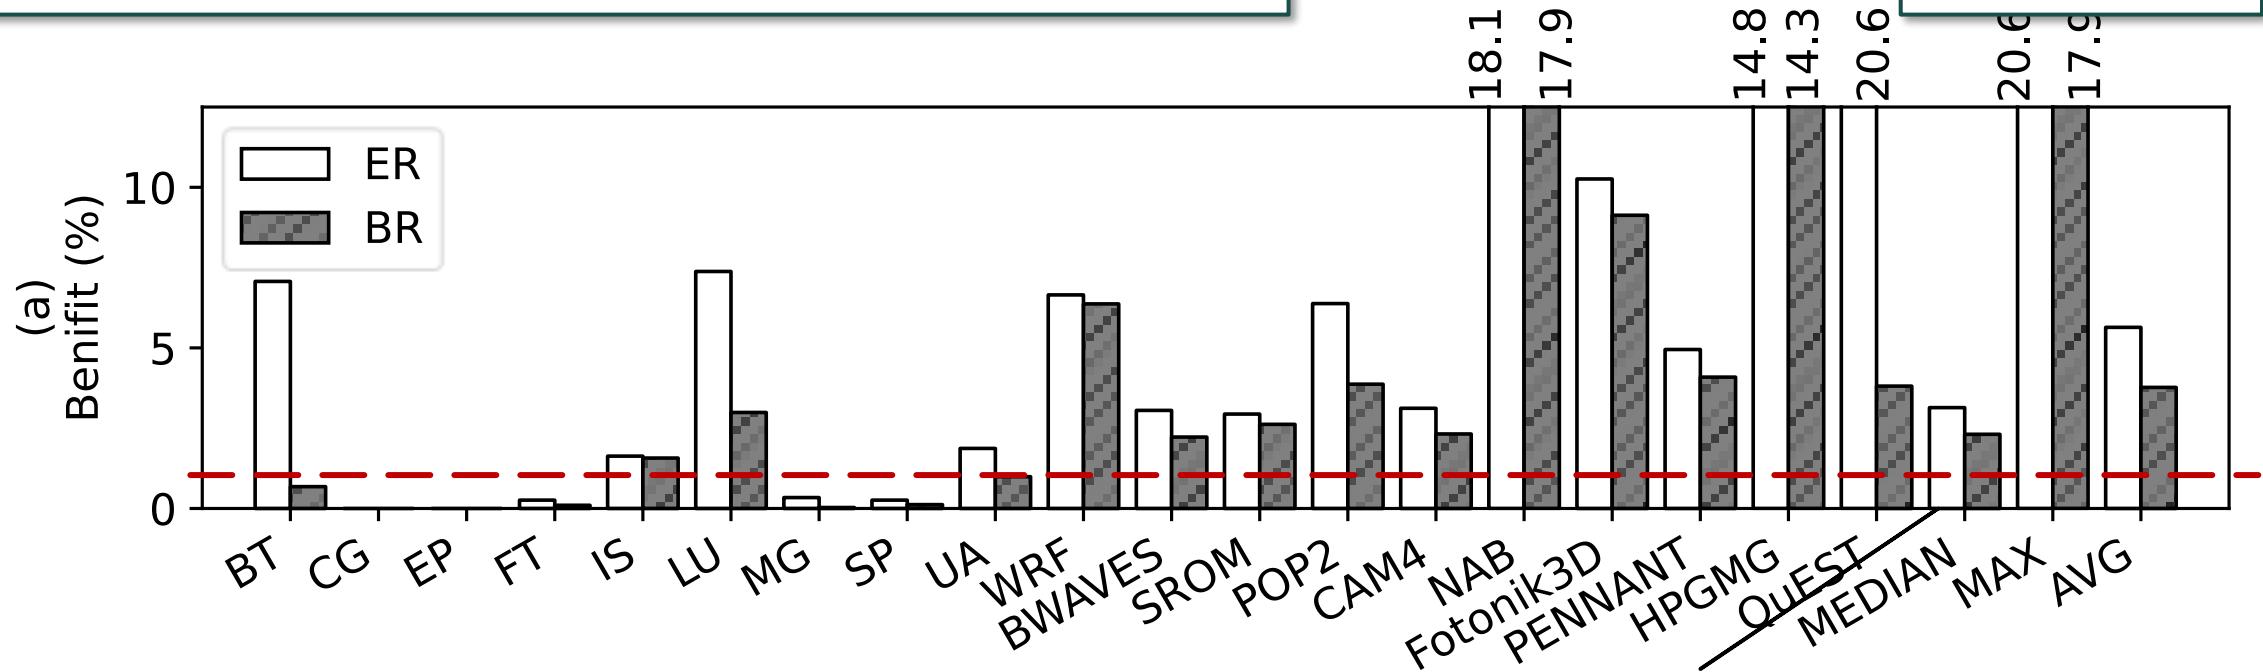
Experimental Setup

- **Server:**
 - 256 GB DDR4 memory & Intel Xeon E52680v4@2.40GHz (14 threads)
 - Linux 5.4.0-77-generic Ubuntu 20.04 LTS
- **Performance potential estimation:** $Latency_{mem} = 50$ cycles
- **Representative programs**
 - NPB 3.4.2 (Class C)
 - SPEC CPU2017 (ref)
 - CORAL2
 - QuEST
- **Compiler:** GCC 9.4.0 –O3 –fopenmp (-g for profiling)

Identifying Trivialities

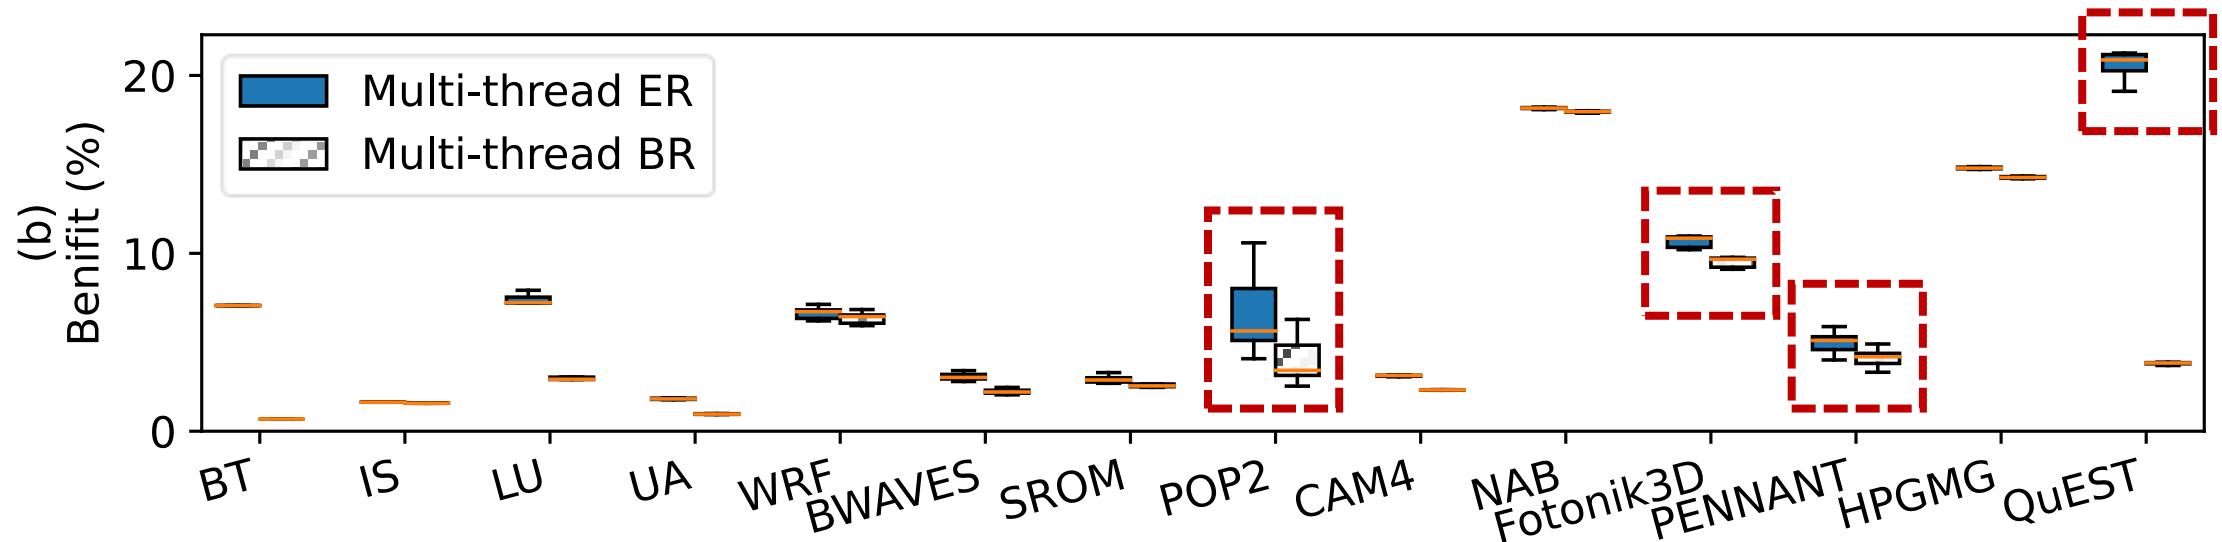
- *ER: Expected Benefit Rate*
- *BR: Branching Benefit Rate*
- **Higher value indicates more opportunity**

AVG:
• *ER 5.64%*
• *BR 3.77%*



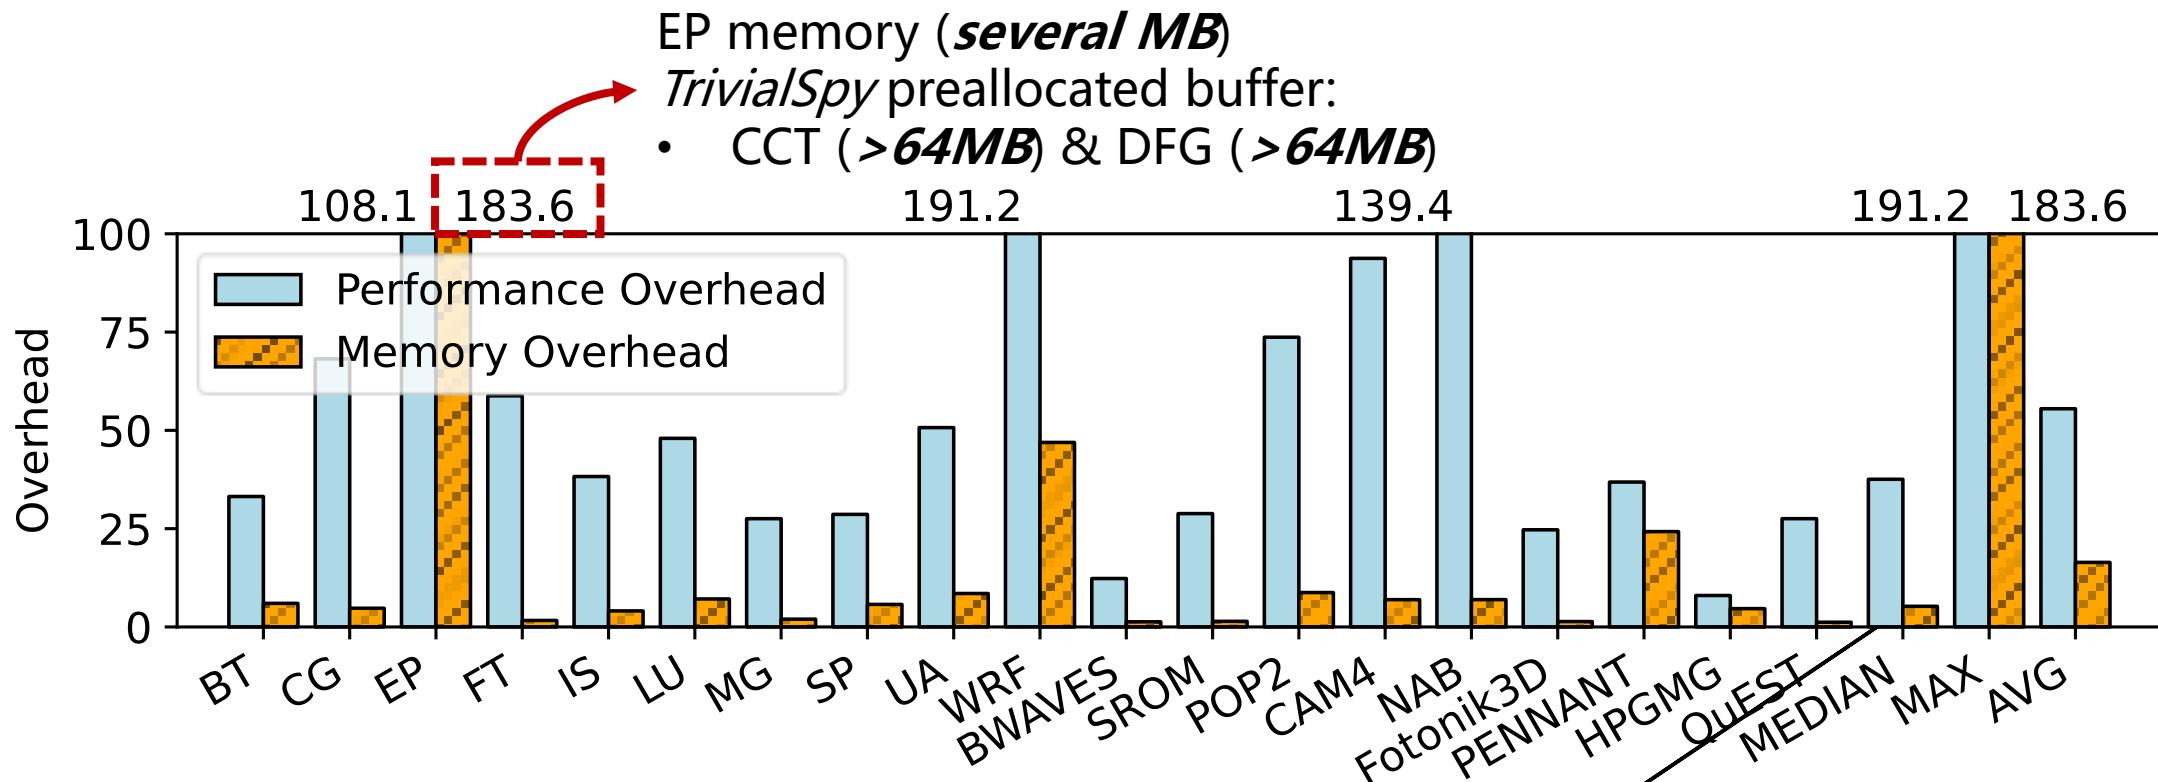
In general, we consider the ***optimization potential is actionable*** when the estimated metric ***ER or BR is larger than 1%***.

Identifying Trivialities



POP2, PENNANT, Fotonik3D and QuEST expose **large variance ($\geq 2.38\%$)** of ER and BR among all threads, which indicates **potential load imbalance** caused by trivial operations.

Overhead



The average performance and memory overhead of all evaluated programs are $55.50\times$ and $16.62\times$, respectively.

The runtime overhead **is similar to well-accepted** binary instrumentation profilers.

Performance Improvement - Overview

- **BO: branch optimization**
- **HTFE: heavy trivial function elimination**
- **LBO: load balance optimization**

Programs	Directed by	Code Line of Triviality	Opt.	Speedup
IS [8]	BB(TCR)	randlc@is.c:369	BO	4.14%±0.80%
LU [8]	BB(RBSR,CTR)	buts@buts.f90:loop(50-69)	BO	4.19%±0.22%
UA [8]	BB(RBSR,CTR)	diffusion@diffuse.f90:140	BO	2.14%±0.74%
WRF [11]	EB(HIR)	psim_unstable@module_sf_sfclayrev.fppized.f90:1098	HTFE	18.51%±0.21%
BWAVES [11]	BB (HIR,RBSR,CTR)	shell@shell_lam.fppized.f:243-270 jacobian@jacobian_lam.fppized.f:94-133	BO	1.05%±0.07%
SROM [11]	BB(TCR,RBSR)	pre_step3d@pre_step3d.fppized.f90:1742	BO	0.90%±0.11%
POP2 [11]	Unbalanced-EB (HIR,CTR)	submeso_flux@mix_submeso.fppized.f90:862 baroclinic_driver@baroclinic.fppized.f90:518	BO + LBO	2.22%±0.21%
CAM4 [11]	BB(HIR,CTR)	cosp_precip_mxratio@cosp_utils.fppized.f90:76	BO	1.47%±0.16%
NAB [11]	EB(HIR)	egb@eff.c:2107	HTFE	9.80%±0.16%
Fotonik3D [11]	BB(CTR)	updateh@update.fppized.f90:loop(189-201)	BO	51.11%±0.3%
PENNANT [1]	Unbalanced-EB (RBSR)	updateh@update.fppized.f90:189	+ LBO	52.09%±0.03%
		QCS::setQCnForce@QCS.cc:270 Hydro::doCycle@Hydro.cc:211	BO + LBO	2.57% ±0.06%
HPGMG [1]	BB(RBSR,CTR)	rebuild_operator_blackbox@rebuild.c:130	BO	1.30%±0.40%
QuEST [26]	BB(CTR)	statevec_controlledCompactUnitaryLocal@QuEST_cpu.c:loop(2241-2257)	BO	14.22%±0.27%
	Unbalanced-EB	statevec_controlledCompactUnitaryLocal@QuEST_cpu.c:2230	+ LBO	23.96%±0.17%

Programs with **high optimization potentials** can obtain performance improvement after optimization.

For more details, please refer to our paper and API docs.

X. You, H. Yang, K. Lei, Z. Luan and D. Qian, "TrivialSpy: Identifying Software Triviality via Fine-grained and Dataflow-based Value Profiling," SC23: International Conference for High Performance Computing, Networking, Storage and Analysis, Denver, CO, USA, 2023, pp. 1-14.

TrivialSpy is open-source: <https://github.com/VClinic/VClinic>

Program	% Benefits			% Detailed Metrics			BR/ER Accuracy	
	ER	BR	HIR	TCR	RBSR	CTR	OP*	% Error†
BT	7.07	0.68	0.36	88.51	5.28	100.0	✗‡	-
CG	0.00	0.00	0.77	0.71	95.20	0.52	✗	-
FT	0.26	0.10	0.00	0.00	2.15	84.77	✗	-
IS	1.63	1.57	0.00	0.00	94.44	0.00	✓	+2.51/+2.57
LU	7.38	2.99	10.24	53.81	78.91	100.0	✓	-3.19/+1.20
MIC	0.03	0.00	0.05	0.00	0.00	0.00	✗	-
SP	0.26	0.12	28.13	44.42	52.17	99.94	✗	-
NWFF	6.67	2.15	38.00	77.00	99.96	100.0	✓	-1.97/-1.16
BWAVES	2.96	2.15	38.42	27.74	41.23	87.73	✓	-2.03/-1.18
SROM	2.94	2.62	30.90	6.68	68.21	63.59	✓	-2.04/-1.72
POP2	6.38	3.87	59.08	27.57	45.70	71.94	✓	-4.16/-1.65
CAM4	3.12	2.32	75.99	13.80	15.12	24.00	✓	-1.65/-0.85
Fotonik3D	10.26	9.13	62.16	14.38	28.64	97.29	✓	+41.83/+42.96
PENNANT	4.95	4.09	5.78	14.73	85.59	56.83	✓	-2.38/-1.52
HPGMG	14.78	14.26	11.55	3.87	98.92	99.87	✓	-13.48/-12.96
QuEST	20.59	3.81	0.00	80.19	77.75	99.74	✓	+3.37/+20.15

Outline

- Introduction & Background
- Understanding Software Triviality
- Dataflow-based Triviality Detection
- Evaluation
- Hands-on Tutorial
 - Installation
 - Case Study – Backprop
 - Case Study – IS Benchmark
 - Case Study – fotonik3d

Installation - TrivialSpy

- **Install with source code:**

- Dependencies: git, gcc>=9, g++, make, cmake>=3.20

- Source Code: `git clone https://github.com/VClinic/VClinic`

- Compilation Instruction: `cd VClinic && ./build.sh`

- Configure the Path: `export DRRUN=`pwd`/build/bin64/drrun`

- **Install in tutorial cluster:**

- Source Code: `cp -r /public/home/buaa_hipo/shared_folder/VClinic ./`

- Compilation Instruction: `cd VClinic && ./build.sh`

- Configure the Path: `export DRRUN=`pwd`/build/bin64/drrun`

- **Instruction to analyze the target program:**

- TrivialSpy: `$DRRUN -t trivialspy -- <EXE> <ARGS>`

Case Study – backprop (~6mins)

- Get the Benchmark and Compile: backprop

- Get the rodinia_3.1:

```
cp -r /public/home/buua_hipo/shared_folder/backprop ./
```

- Compile:

```
cd /backprop  
vim Makefile # add the -g to all flags  
make  
cp slurm-template.sh trivial-backprop.sh  
vim trivial-backprop.sh  
sh trivial-backprop.sh
```

Modify the Job Name

```
NTASKS=1  
#!/bin/bash  
JOBNAME="vcclinic-trivialspy-backprop"  
DRRUN=" `pwd` /VClinic/build/bin64/drrun"  
TOOL="trivialspy"  
# TARGET command for profiling: CMD="<EXE> <ARGS>"  
CMD="`pwd`/backprop "  
mkdir -p log  
  
echo "START $JOBNAME WITH NTASK=$NTASKS "  
nowdate=$(date +%Y_%m_%d_%H_%M_%S)  
echo $nowdate  
sbatch << END
```

Modify the EXE PATH

```
#!/bin/bash  
#SBATCH -J $JOBNAME  
#SBATCH -o log/$JOBNAME-$NTASKS-%j-$nowdate.log  
#SBATCH -e log/$JOBNAME-$NTASKS-%j-$nowdate.err  
#SBATCH -p test  
#SBATCH --cpus-per-task=16  
#SBATCH --ntasks-per-node=1  
#SBATCH -n $NTASKS  
# Your SCRTPT commands  
time $CMD 6553600  
time $DRRUN -t $TOOL -- $CMD 65536  
END
```

Unoptimized time

Profiling time

Case Study – backprop

- Resulting files are generated in the x86-<host>-<PID>-trivialspray folder
 - trivialspray.log is the summary report for invalid operation detection metrics
 - thread-<id>.log contains the invalid operation detection reports for individual threads.
 - Summary reports: backprop has an expected benifit of 1.300 and a branch benifit of 0.866.

```
Running: /public/home/csjt0800/xzb/VClinic/build/bin64/..../clients/lib64/release/libtrivialspray.so
[TRIVIALSPY INFO] Profiling with value tracing
[TRIVIALSPY INFO] Thread Private is disabled.
[TRIVIALSPY INFO] Soft Approximation is disabled
[TRIVIALSPY INFO] Hard Approximation is disabled

----- [Thread=1] Dumping Dataflow-aware Trivial Inefficiency Overview Report -----
```

Total Speculate Benifit: 1.300 (1570475 benifit / 120837540 total cost)

Total Benifit: 0.866 (1046821 benifit / 120837540 total cost)

Total Heavy Instruction: 0.010 (153 / 1570475 SB)

Total Trivial Chain: 49.996 (785176 / 1570475 SB)

Total Redundant Backward Slice: 0.004 (58 / 1570475 SB)

Total Absorbing Breakpoints: 99.979 (1570140 / 1570475 SB)

trivialspray.log

Case Study – backprop

- Optimization Guide: the xmm0, xmm1 of backprop.c:369 leading to chained trivial operation

```
+++ DFG Summary Info +++
===== DFGLog from Thread 2761 =====
exe count: 130024
total cost: 433
benifit: 12
TC branch cost: 4
heavy cost: 0
chained cost: 6
backward slice cost: 0
Trivial Condition Num: 2

+++++ Singular Trivial Condition(s):
<8> <4, 5>: opnd=xmm0, val=ZERO, isSingular=1
<20> <11, 12>: opnd=xmm1, val=ZERO, isSingular=1
=> detailed node info:
[4] cvtss2sd xmm0, ..... backprop.c:323]
[5] mulsd xmm0, xmm2 ..... backprop.c:323] [D]
[6] cvtss2sd xmm1, dword ptr [r9+rax*4] ..... backprop.c:323] [B]
[10] pxor xmm1, xmm1 ..... backprop.c:323]
[11] cvtss2sd xmm1, dword ptr [rdx] ..... backprop.c:323]
[12] mulsd xmm1, xmm2 ..... backprop.c:323] [D]
[13] addsd xmm0, xmm1 ..... backprop.c:323] [D][P]

=> detailed edge info:
<4> <-1, 2>: opnd=xmm1, val=UNKNOWN, isSingular=0
<5> <-1, 3>: opnd=qword ptr [r10+rax*8], val=UNKNOWN, isSingular=0
<11> <1, 7>: opnd=rdx, val=UNKNOWN, isSingular=0
<12> <-1, 8>: opnd=rsi, val=UNKNOWN, isSingular=0
```

thread-1.log

```
#ifdef OPEN
    omp_set_num_threads(NUM_THREAD);
    #pragma omp parallel for \
        shared(oldw, w, delta) \
        private(j, k, new_dw) \
        firstprivate(ndelta, nly)
#endif
    for (j = 1; j <= ndelta; j++) {
        for (k = 0; k <= nly; k++) {
            new_dw = ((ETA * delta[j] * ly[k]) + (MOMENTUM * oldw[k][j]));
            w[k][j] += new_dw;
            oldw[k][j] = new_dw;
        }
    }
}
```

backprop.c

In many cases, both **delta[j]** and **oldw[k][j]** will be zero, resulting in a large number of trivial operations for updating **new_dw**, **w[k][j]**, and **oldw[k][j]**.

new_dw = 0 if delta[j]==0 && oldw[k][j]==0.

Case Study – backprop

- Optimization Guide: add the branch code.

```
#ifdef OPEN
    omp_set_num_threads(NUM_THREAD);
    #pragma omp parallel for \
        shared(oldw, w, delta) \
        private(j, k, new_dw) \
        firstprivate(ndelta, nly)
#endif
    for (j = 1; j <= ndelta; j++) {
        for (k = 0; k <= nly; k++) {
            new_dw = ((ETA * delta[j] * ly[k]) + (MOMENTUM * oldw[k][j]));
            w[k][j] += new_aw;
            oldw[k][j] = new_dw;
        }
    }
}
```

backprop.c

In many cases, both **delta[j]** and **oldw[k][j]** will be zero, resulting in a large number of trivial operations for **new_dw**.

vim backprop/backprop.c

```
#ifdef OPEN
    omp_set_num_threads(NUM_THREAD);
    #pragma omp parallel for \
        shared(oldw, w, delta) \
        private(j, k, new_dw) \
        firstprivate(ndelta, nly)
#endif
    for (j = 1; j <= ndelta; j++) {
        for (k = 0; k <= nly; k++) {
            if (delta[j]==0 && oldw[k][j]==0){
                new_dw=0;
            } else{
                new_dw = ((ETA * delta[j] * ly[k]) + (MOMENTUM * oldw[k][j]));
                w[k][j] += new_dw;
                oldw[k][j] = new_dw;
            }
        }
    }
}
```

backprop.c

new_dw = 0 if delta[j]==0 && oldw[k][j]==0.

make #recompilation

Improve 20.6% performance

Case Study – IS Benchmark (~6mins, Optional)

- Get the Benchmark and Compile: IS - Integer Sort, random memory access
 - Get the NPB Benchmark:

```
cp -r /public/home/buua_hipo/shared_folder/NPB-3.4.2-OMP ./
```

- Compile:

```
cd NPB3.4-OMP/  
cp config/make.def.template config/make.def  
vim config/make.def # add the -g to all flags  
make IS CLASS=C && cd ..  
cp slurm-template.sh trivial-is.sh  
vim trivial-is.sh && sh trivial-is.sh
```

```
#-----  
# Global *compile time* flags for Fortran programs  
#-----  
FFLAGS = -O3 -fopenmp -g  
  
#-----  
# _FILE_OFFSET_BITS=64  
# _LARGEFILE64_SOURCE - are standard compiler flags which  
# files larger than 2GB.  
#-----  
CFLAGS = -O3 -fopenmp -g
```

Modify the Job Name

```
NTASKS=1  
#!/bin/bash  
JOBNAME="vcclinic-trivialspy-is"  
DRRUN=" `pwd`/VClinic/build/bin64/drrun"  
TOOL="trivialspy"  
# TARGET command for profiling: CMD="<EXE> <ARGS>"  
CMD="`pwd`/NPB3.4-OMP/bin/is.C.x"  
mkdir -p log  
  
echo "START $JOBNAME WITH NTASK=$NTASKS "  
nowdate=$(date +%Y_%m_%d_%H_%M_%S)  
echo $nowdate  
sbatch << END
```

Modify the EXE PATH

```
#!/bin/bash  
#SBATCH -J $JOBNAME  
#SBATCH -o log/$JOBNAME-$NTASKS-%j-$nowdate.log  
#SBATCH -e log/$JOBNAME-$NTASKS-%j-$nowdate.err  
#SBATCH -p test  
#SBATCH --cpus-per-task=16  
#SBATCH --ntasks-per-node=1  
#SBATCH -n $NTASKS  
# Your SCRIPT commands  
cd `pwd`/PENNANT/build  
time -- $CMD  
time $DRRUN -t $TOOL -- $CMD  
END
```

Unoptimized time

Unoptimized time

Case Study – IS Benchmark (Optional)

- Analyze the IS with TrivialSpy: `$DRRUN -t trivialspy -- ./bin/is.C.x`
- Resulting files are generated in the x86-<host>-<PID>-trivialspy folder
 - trivialspy.log is the summary report for invalid operation detection metrics
 - thread-<id>.log contains the invalid operation detection reports for individual threads.
 - Summary reports: IS has an expected benifit of 25.003 speculate benifit.

==== Overall Triviality Metric ====
trivialspy.log
----- [Thread=2] Dumping Dataflow-aware Trivial Inefficiency Overview Report -----

Total Speculate Benifit: 0.089 (74528504 benifit / 84005563936 total cost)
Total Benifit: 0.030 (24842849 benifit / 84005563936 total cost)
Total Heavy Instruction: 0.000 (102 / 74528504 SB)
Total Trivial Chain: 0.000 (217 / 74528504 SB)
Total Redundant Backward Slice: 0.000 (279 / 74528504 SB)
Total Absorbing Breakpoints: 0.001 (402 / 74528504 SB)

----- [Thread=2] Dumping Dataflow-aware Trivial Inefficiency Report -----
----- Trivial Hotspots ordered by BB -----

Total Speculate Benifit: 0.089 (74528504 benifit / 84005563936 total cost)
Total Benifit: 0.030 (24842849 benifit / 84005563936 total cost)
Total Heavy Instruction: 0.000 (102 / 74528504 SB)
Total Trivial Chain: 0.000 (217 / 74528504 SB)
Total Redundant Backward Slice: 0.000 (279 / 74528504 SB)
Total Absorbing Breakpoints: 0.001 (402 / 74528504 SB)

Benifit: 25.003 (6211538 local benifit / 24842849 total benifit)
Importance: 0.007 (6211538 benifit / 84005563936 total cost)

Speculate Benifit: 25.003 (18634614 local SB / 74528504 SB)
Importance: 0.022 (18634614 benifit / 84005563936 total cost)

Case Study – IS Benchmark (Optional)

- Optimization Guide: the 0 value xmm3 of is.c:369 leading to chained trivial operation
 - Eliminate the propagation of invalid calculations caused by zero values resulting from converting.

```
+++ DFG Summary Info +++
===== DFGLog from Thread 31647 =====
exe count: 6210686
total cost: 278
benifit: 3
TC branch cost: 2
heavy cost: 0
chained cost: 0
backward slice cost: 0
Trivial Condition Num: 1
+++++ Singular Trivial Condition(s):
<55> <34, 35>: opnd=xmm3, val=ZERO, isSingular=1
--> detailed node info.
[33] pxor xmm3, xmm3
@randlc[/public/home/buaa_hipo/app/NPB3.4.2/NPB3.4-OMP/IS/is.c:368]
[34] cvtsi2sd xmm3, eax
@randlc[/public/home/buaa_hipo/app/NPB3.4.2/NPB3.4-OMP/IS/is.c:368]
[35] mulsd xmm1, xmm3
@randlc[/public/home/buaa_hipo/app/NPB3.4.2/NPB3.4-OMP/IS/is.c:369] [D]
[36] subsd xmm2, xmm1
@randlc[/public/home/buaa_hipo/app/NPB3.4.2/NPB3.4-OMP/IS/is.c:369] [P]
...
vim NPB3.4-OMP/IS/is.c
```

```
j = R23 * T1;
T2 = j;
Z = T1 - T23 * T2;
T3 = T23 * Z + A2 * X2;
j = R46 * T3;
T4 = j;
*X = T3 - T46 * T4;
return(R46 * *X);

j = R23 * T1;
T2 = j;
Z = T1 - T23 * T2;
T3 = T23 * Z + A2 * X2;
j = R46 * T3;
T4 = j;
if(T4==0) {
    *X = T3;
    return R46*T3;
}
*X = T3 - T46 * T4;
return(R46 * *X);
```

Improve 5% performance

Case Study – fotonik3d

- Get the Benchmark and Compile: spec cpu2017/fotonik3d

- Get the fotonik3d:

```
cp -r /public/home/buua_hipo/shared_folder/649.fotonik3d ./
```

- Compile:

```
cd 649.fotonik3d  
make && cd ..  
cp slurm-template.sh trivial-fotonik.sh  
vim trivial-fotonik.sh  
sh trivial-fotonik.sh
```

Modify the Job Name

```
NTASKS=1  
#!/bin/bash  
JOBNAME="vcclinic-trivalspy-fotonik"  
DRRUN=`pwd`/VClinic/build/bin64/drrun"  
# tool name example: zerospy, trivalspy  
TOOL="trivalspy"  
# TARGET command for profiling: CMD=<EXE> <ARGS>"  
CMD="fotonik3d s base.mytest-m64"  
mkdir -p log  
echo "START $JOBNAME WITH NTASK=$NTASKS "  
nowdate=$(date +%Y_%m_%d_%H_%M_%S)  
echo $nowdate  
sbatch << END
```

Modify the EXE PATH

```
#!/bin/bash  
#SBATCH -J $JOBNAME  
#SBATCH -o log/$JOBNAME-$NTASKS-%j-$nowdate.log  
#SBATCH -e log/$JOBNAME-$NTASKS-%j-$nowdate.err  
#SBATCH -p test  
#SBATCH --cpus-per-task=16  
#SBATCH --ntasks-per-node=1  
#SBATCH -n $NTASKS  
# Your SCRIPT commands  
cd 649.fotonik3d  
time $CMD  
time $DRRUN -t $TOOL -- $CMD  
END
```

Unoptimized time

Profiling time

Case Study – fotonik3d

```
===== DFGLog from Thread 16547 =====
exe count: 8363306
total cost: 3499
benifit: 1425
TC branch cost: 180
heavy cost: 960
chained cost: 312
backward slice cost: 312
Trivial Condition Num: 15
++++++ Singular Trivial Condition(s):
<0> <-1, 0>: opnd=qword ptr [r15+rax*8], val=ZERO, isSingular=1
<9> <6, 7>: opnd=xmm2, val=ZERO, isSingular=1
<11> <7, 8>: opnd=xmm2, val=ZERO, isSingular=1
<14> <8, 9>: opnd=xmm1, val=ZERO, isSingular=1
<28> <16, 17>: opnd=xmm1, val=ZERO, isSingular=1
<38> <-1, 23>: opnd=qword ptr [rbx+rax*8], val=ZERO, isSingular=1
<46> <28, 29>: opnd=xmm2, val=ZERO, isSingular=1
<48> <29, 30>: opnd=xmm2, val=ZERO, isSingular=1
<51> <30, 31>: opnd=xmm1, val=ZERO, isSingular=1
<65> <37, 39>: opnd=xmm1, val=ZERO, isSingular=1
<75> <-1, 45>: opnd=qword ptr [r9+rax*8], val=ZERO, isSingular=1
<83> <50, 51>: opnd=xmm2, val=ZERO, isSingular=1
<85> <51, 52>: opnd=xmm2, val=ZERO, isSingular=1
<88> <52, 53>: opnd=xmm1, val=ZERO, isSingular=1
<101> <59, 60>: opnd=xmm1, val=ZERO, isSingular=1
==> detailed node info:
[0] vmovsd xmm0, qword ptr [r15+rax*8] @_upml_mod_MOD_upml_updateh._omp_fn.0[/public/home/csjt0800/lkl/.loc
uild/build_base_mytest-m64.0000/UPML.fppized.f90:1490] [D][P]
[1] mov rsi, qword ptr [rsp+0x10] @_upml_mod_MOD_upml_updateh._omp_fn.0[/public/home/csjt0800/lkl/.loc
ld/build_base_mytest-m64.0000/UPML.fppized.f90:1494]
[2] vmovsd xmm1, qword ptr [rsi+rax*8] @_upml_mod_MOD_upml_updateh._omp_fn.0[/public/home/csjt0800/lkl/.loc
uild/build_base_mytest-m64.0000/UPML.fppized.f90:1494] [D][B]
[3] vsubsd xmm1, xmm1, qword ptr [rdx+0x08] @_upml_mod_MOD_upml_updateh._omp_fn.0[/public/home/csjt0800/l
d_s/build/build_base_mytest-m64.0000/UPML.fppized.f90:1494] [D][B]
[4] vmovsd xmm2, qword ptr [rcx+0x08] @_upml_mod_MOD_upml_updateh._omp_fn.0[/public/home/csjt0800/lkl/.loc
ild/build_base_mytest-m64.0000/UPML.fppized.f90:1494]
[5] mov rdi, qword ptr [rsp+0x30] @_upml_mod_MOD_upml_updateh._omp_fn.0[/public/home/csjt0800/lkl/.loc
ld/build_base_mytest-m64.0000/UPML.fppized.f90:1494]
[6] vsubsd xmm2, xmm2, qword ptr [rdi+rax*8] @_upml_mod_MOD_upml_updateh._omp_fn.0[/public/home/csjt0800/
3d_s/build/build_base_mytest-m64.0000/UPML.fppized.f90:1494]
[7] vmulsd xmm2, xmm2, <rel> qword ptr [0x00000000006673e8] @_upml_mod_MOD_upml_updateh._omp_fn.0[/public/
CPU/649.fotonik3d_s/build/build_base_mytest-m64.0000/UPML.fppized.f90:1494] [D][H]
[8] vfmadd132sd xmm1, xmm2, <rel> qword ptr [0x00000000006673d8] @_upml_mod_MOD_upml_updateh._omp_fn.0[/p
spec/CPU/649.fotonik3d_s/build/build_base_mytest-m64.0000/UPML.fppized.f90:1494] [D][H]
[9] vmulsd xmm1, xmm1, qword ptr [r12] @_upml_mod_MOD_upml_updateh._omp_fn.0[/public/home/csjt0800/lkl/.loc
uild/build_base_mytest-m64.0000/UPML.fppized.f90:1494] [D][H]
[10] mov rsi, qword ptr [rsp] @_upml_mod_MOD_upml_updateh._omp_fn.0[/public/home/csjt0800/lkl/.local/cp
ild_base_mytest-m64.0000/UPML.fppized.f90:1494]
[11] vfmadd231sd xmm1, xmm0, qword ptr [rsi] @_upml_mod_MOD_upml_updateh._omp_fn.0[/public/home/csjt0800/l
d_s/build/build_base_mytest-m64.0000/UPML.fppized.f90:1494]
```

```
!$OMP DO PRIVATE(i,j,k)
do k=zstart,0
  do j=1,ny
    do i=1,nx
      Bxold = Bx_klow(i,j,k)
      Bx_klow(i,j,k) = ayh(j) * Bx_klow(i,j,k) +
        byh(j) * ((Ey(i,j,k+1)-Ey(i,j,k)) * dzinv +
        (Ez(i,j,k)-Ez(i,j+1,k)) * dyinv)
      Hx(i,j,k) = azh(k) * Hx(i,j,k) +
        bzh(k) * (cxe(i)*Bx_klow(i,j,k) - fxe(i)*Bxold) * muinv
      !-----
      Byold = By_klow(i,j,k)
      By_klow(i,j,k) = azh(k) * By_klow(i,j,k) +
        bzh(k) * ((Ex(i,j+1,k)-Ex(i,j,k)) * dxinv +
        (Ey(i,j,k)-Ey(i,j+1,k)) * dzinv)
      Hy(i,j,k) = axh(i) * Hy(i,j,k) +
        bxh(i) * (cye(j)*By_klow(i,j,k) - fye(j)*Byold) * muinv
      !-----
      Bzold = Bz_klow(i,j,k)
      Bz_klow(i,j,k) = axh(i) * Bz_klow(i,j,k) +
        bxh(i) * ((Ex(i,j+1,k)-Ex(i,j,k)) * dyinv +
        (Ey(i,j,k)-Ey(i,j+1,k)) * dxinv)
      Hz(i,j,k) = ayh(j) * Hz(i,j,k) +
        byh(j) * (cze(k)*Bz_klow(i,j,k) - fze(k)*Bzold) * muinv
      end do
    end do
  end do
```

Thanks! Q&A